PsicoAyudaVen — 12-month roadmap

Status: active planning document (v1, July 2026) **Horizon:** July 2026 → July 2027 **Owner:** tech lead + clinical advisor

This roadmap consolidates the strategic plan for the next 12 months. It is built on the current production state (v1.23.0) and reflects decisions made across product, architecture, team, and funding. It is the artifact to share with funders, partners, and new team members.

For implementation gotchas, see `AGENTS.md`. For shipped history, see `CHANGELOG.md`.

---

1. Executive summary

PsicoAyudaVen is a disaster-response psychological-support platform connecting people in Venezuela with verified psychologists, evolving toward a volunteer *acompañamiento* model (cross-border, non-therapy). In the next 12 months it grows from a software-engineer-led, clinically-guided PWA into a small-team operation shipping **native iOS + Android apps with embedded video sessions**, a **professional AI copilot**, and **scheduling for ongoing therapy relationships** — while keeping the instant crisis-triage path that is the platform's reason for existing.

Year-1 envelope: ~$220,000–400,000 (team + platform), pursued via funding collectives. Today the platform is sustained by a single developer; the code is open source (MIT) on GitHub.

---

2. Where we are (July 2026)

Dimension	Current state
Team	1 software engineer (sole supporter — funds, builds, and runs the platform); clinical guidelines set by an advisory group of mental-health medical professionals

Platform cost	~$0–30 / month (Cloudflare free tiers + small Sentry)
Surface	PWA only — installs on Android, no App Store / Play Store presence
Real-time	None — contact is WhatsApp outbound + pre-recorded audio stories
Codebase	One TanStack Start app (React 19 + Cloudflare Workers + D1 + R2 + Analytics Engine)
Code	**Open source** (MIT license), public on GitHub (`github.com/enBonnet/psico-support`)
Model	Verified-pro directory operational (v1.0, 2026-06-28); volunteer *acompañamiento* model drafted (`/voluntariado` , `/privacidad`) but **not legally validated**, no volunteer registration flow
Release cadence	~25 releases in ~25 days (v1.0 → v1.23.0)

Governance today: the platform is sustained solely by its lead software engineer, who funds, builds, and operates it. An advisory group of mental-health medical professionals defines the clinical guidelines, protocols, and professional standards that therapists and volunteers on the platform must follow — the engineer executes delivery; the clinical group sets the standard of care. This split already exists and is the foundation the Year-1 team expansion builds on — the clinical-advisor role in §12 formalizes and extends it, not creates it. The codebase is open source under the MIT license, born at the Build4Venezuela open hackathon (a free, no-prize community event — **not** a founder, funder, or ongoing institutional relationship).

Annual run-rate today: ~$200–400 (domain + minimal services).

Metrics since launch (2026-06-28):

Metric	Value	Source
Unique visitors	4,140	Cloudflare Web Analytics
Unique visitors/day (current)	~300	Cloudflare Web Analytics
Engaged actors (≥1 tracked event)	651 (~16% of visitors)	Analytics Engine
Total tracked events	2,538	Analytics Engine

WhatsApp contacts (all CTAs)	~109 — help-now 64, directory card 27, `/ahora` 10, random 8	Analytics Engine
PWA installs	6	Analytics Engine

Reach vs engagement: ~~4,140 visitors reached the site, but only ~651 (~~16%) triggered a tracked action and ~~109 (~~2.6%) opened a WhatsApp conversation. Engaged actors peaked at 212 on launch day (Jul 2) and now run ~30–60/day against ~300 raw visitors/day. The help-now CTA is the single biggest converter (64 of 109 contacts). This gap is the quantitative case for the funnel-rebuild work in Q4 — the platform has reach; converting that reach into help-seeker contacts is the open product problem.

---

3. Where we want to be (July 2027)

Dimension	Target state
Team	4–6 people (3 full-time eng + clinical advisor + design + community)
Surface	Web (PWA, public reach) + **native iOS + Android apps** in the stores
Real-time	**Embedded video sessions** (LiveKit) — the primary professional channel
Chat	In-app first-aid chat (Cloudflare Durable Objects) + WhatsApp (demoted to first-aid)
AI	**Pro-facing copilot** (follow-up assist, recommendations, risk patterns). Help-seeker never interacts with AI.
Scheduling	Video-session booking with calendars, cases, longitudinal therapy tracking
Model	Volunteer model legally validated and operational alongside the verified-pro directory
Funding	Sustainable source identified and committed
Measurement	Outcome metrics (not just funnel)

---

4. The platform contains two coexisting products

This split is load-bearing and must not be blurred in the UI.

Product	Goal	Session	Friction	Commitment	Identity	Channel
Crisis triage (existing)	**Restore stability** — PFA, not therapy	~15 min, on-demand	Zero — no account, no booking	Single contact	Anonymous OK	Chat / WhatsApp / on-demand video
Ongoing therapy (new)	Therapeutic work over time	Scheduled, longer	Booking, account, consent	Multi-session relationship	Identified	Scheduled video sessions

Rule: a person in crisis must never hit a "create account to talk now" wall. The crisis path stays frictionless; the therapy path is a committed, identified relationship.

Clinical boundary (load-bearing): the crisis path delivers **psychological first aid (PFA) — stabilization, not therapy.** Its sole goal is to restore the person's stability in the moment (a ~15-minute support session) and refer onward if needed. It is **not** psychotherapy and **not** regulated professional practice, which is precisely why trained volunteers can deliver it (the volunteer *acompañamiento* model serves this path). The therapy path is actual therapeutic work and remains the exclusive domain of licensed professionals. These goals must not blur: a crisis session never becomes ad-hoc therapy, and a therapy engagement never collapses into a 15-minute patch.

---

5. Channel hierarchy

```
Help-seeker arrives (web or native app)
  │
  ├─ CRISIS PATH (~15-min stabilization — PFA, not therapy)
  │    ├─ In-app first-aid chat (Durable Objects)
  │    ├─ WhatsApp (ws) — first aid only
  │    └─ On-demand video (~15 min — restore stability, refer onward)
  │
  └─ THERAPY PATH (scheduled therapeutic work — licensed pros only)
       └─ Scheduled video session (LiveKit)

Professional side (all sessions, both paths)
  ├─ Case management (cases + sessions + follow-ups)
  ├─ AI copilot (pro-facing, reviewed, never user-facing)
  └─ Scheduling (therapy-path video sessions only)
```

Video serves both paths with different goals: crisis path = ~15-minute on-demand stabilization (PFA); therapy path = scheduled therapeutic sessions. Chat and WhatsApp are never scheduled — they stay instant first-aid. This confines scheduling complexity to the therapy path's video surface.

---

6. Architecture — backend owns the logic

The single most important architectural principle for a small team:

> **All business logic lives in the backend (Cloudflare Workers). The three presentation layers (web, iOS, Android) are thin renderers that call the same API. No logic is duplicated across platforms.**

```
┌──────────────────────────────────────────────────────────────┐
│                     PRESENTATION (thin)                        │
│   Web PWA (TanStack Start)   │   Native apps (React Native)   │
│   - public reach + SEO       │   - iOS + Android, one RN      │
│   - directory, /recursos,    │     codebase                   │
│     /apoyo, legal pages      │   - the primary product:       │
│   - low-end-device fallback  │     auth, panel, cases,        │
│                              │     video, chat, scheduling    │
└──────────────┬───────────────┴────────────────┬───────────────┘
               │                                │
               └───────────────┬────────────────┘
                               │  shared API contract
                               │  (Zod schemas, typed server fns)
                               ▼
┌───────────────────────────────────────────────────────────────┐
│                  BACKEND (all logic lives here)                 │
│                  Cloudflare Workers                              │
│   ├─ D1 (SQLite) — directory, users, cases, sessions, slots     │
│   ├─ R2 — avatars, certificates, audio stories                  │
│   ├─ Durable Objects — real-time first-aid chat                  │
│   ├─ Analytics Engine — product telemetry (write-only)          │
│   ├─ Email Service — transactional + session reminders           │
│   ├─ LiveKit (external) — video sessions (room creation,        │
│   │                        tokens, webhook handling)            │
│   └─ AI providers — Workers AI (in-region) + LLM API (DPA)      │
```

```
│                       for the pro copilot                         │
└───────────────────────────────────────────────────────────────┘
```

Implications:

- Validation, scheduling rules, conflict checks, risk escalation, AI orchestration, video token minting — **all server-side.**
- Clients render state and call mutations. Nothing smarter.
- The shared schema/type layer (Zod) is the contract. A change to a schema is the single source of truth across all three clients.
- **One mobile engineer can ship both iOS and Android** because the RN app is a thin renderer. This is the "more with less" lever.

---

7. AI strategy — pro-facing copilot only

Principle: the help-seeker never interacts with AI. The "personas reales, no bots ni inteligencia artificial" promise (landing, v1.13.1) holds without rewrite. AI is a tool the professional uses on the back end; everything it produces is reviewed by a human before reaching any record or person.

Copilot features

Feature	What it does	Phase
Follow-up note assistance	Suggests structure, flags missing elements, completes templates from pro notes	Q1 2027
Resource recommendations	Suggests `/recursos`, referrals, worksheets based on session tags	Q1 2027
Caseload risk pattern detection	Flags escalating risk / missed sessions across a pro's open cases	Q2 2027
Session handoff summary	Summarizes the preceding first-aid chat for the incoming pro	Q1 2027
Treatment plan drafting	Suggests goals/milestones from case context (pro reviews)	Q2 2027
Outcome-over-time signals	Surfaces progress signals across sessions	Q2 2027

Provider strategy (two-tier, data-handling tied to legal review)

- **Cloudflare Workers AI** (in-region) for cheap/fast calls: tagging, classification, routing flags. Data stays in Cloudflare — no DPA needed.
- **LLM API** (Anthropic/OpenAI) for clinical-grade generation (drafting, summaries). Requires a **Data Processing Agreement + explicit consent**. Decision deferred to Q3 legal review.

Off the table (non-negotiable):

- AI chatbot *as therapist*.
- AI diagnosis.
- AI without a human escalation path.
- Any AI the help-seeker interacts with directly.

---

8. Data & retention

Data class	Retention
Anonymous crisis-triage contacts (no account)	12 months after last contact (existing draft)
Ongoing therapy cases	**Kept while the case is active.** The 12-month window starts when the case concludes (`status → concluded` or `referred_out`).
Professional verification documents	Kept while the pro is verified; removed on request / removal
Session video	**Never recorded.** LiveKit sessions are live-only.
Analytics	Aggregate-only, immutable column contract (gotcha #10)

Privacy rule: longitudinal case data is far more sensitive than single crisis contacts. External LLM use on session notes requires strong DPA + explicit consent — default to Workers AI (in-region) until legal clears external providers.

---

9. Native apps — full product, React Native, thin

- **One React Native codebase** → iOS + Android.
- Full feature parity with the web product (the app *is* the product).
- Thin renderer — all logic via the shared API (see §6).
- **Web (PWA) stays** as the public reach layer: landing, public directory, pro profiles (SSR for OG/JSON-LD per gotcha #6), `/recursos`, `/apoyo`, legal pages, and a low-end-device fallback.
- New interactive features (video, chat, scheduling, AI copilot) target native first; web follows or stays public-only.

Mobile engineering strategy: start with **one senior mobile engineer**, add a mid/contractor in Q1 2027 if scope demands. Because the app is thin, one engineer can carry both platforms.

---

10. Scheduling, cases, long-term therapy

What exists

- `availability_mode` (always/scheduled/inactive) + recurring weekly blocks
  - IANA timezone, availability derived at view time (v1.13.0, migration 0014).
- `follow_ups` table — flat, one record per contact, owner-scoped (v1.13.0, migration 0013).

What's added (additive migrations, non-breaking)

```
cases (NEW)
  id, professional_id, help_seeker_user_id (nullable),
  status (active|paused|concluded|referred_out),
  treatment_goals (JSON text), started_at, concluded_at

sessions (NEW)
  id, case_id (nullable — crisis contacts stay case-less),
  professional_id, scheduled_at, duration_min,
  status (scheduled|completed|cancelled|no_show),
  video_room_id, risk_level, action_taken, notes

slots (NEW)
  id, professional_id, start_at, end_at,
  status (open|booked|blocked), booked_by_user_id, case_id
```

```
follow_ups (EXISTING — keep, extend)
  + optional case_id (nullable, backfills to null)
```

Build vs buy (scheduling)

- **v1: build minimal on D1** (slots + bookings + conflict checks). The existing `availability_schedule` already models recurring blocks — extend it to materialize bookable slots.
- **v2 (Q2 2027, only if pros demand):** Google Calendar API sync (free).
- Calendly (per-seat cost) and self-hosted Cal.com (ops burden, Node not Workers-native) are **not** adopted.

Scope discipline — this is NOT an EHR

Build the minimum that supports *acompañamiento* continuity. Avoid: structured diagnosis fields, insurance/billing, prescription tracking, medical-grade audit logging. Those are a different regulated company.

---

11. Beyond Year 1 — subscriptions & specialist compensation (Preply-for-therapy)

This is the goal AFTER the 12-month roadmap, not a Year-1 deliverable. Nothing in this section is built during Year 1; it is the endgame that the Year-1 work (team, scheduling, cases, video, native apps, verified specialist pool) is the *foundation* for. Year 1 remains fully grant/collective-funded (see §15); subscriptions begin contributing in Year 2.

The platform will operate a **two-tier service model**, and only one tier is paid:

Tier	Path	Cost	Goal
Free	Crisis triage (~15-min PFA stabilization)	$0 — the mission	Restore stability, refer onward
Subscription	Ongoing therapy (scheduled video sessions)	Monthly fee → quota of sessions	Therapeutic work with licensed specialists

How the subscription works (Preply-style marketplace)

- A user subscribes **monthly** and receives a **quota of therapy sessions** for that month (e.g. a tier might include 4 scheduled sessions/month).
- The user is matched with licensed specialists from the verified directory.
- **Each delivered session pays the specialist.** The platform collects the subscription, disburses per-session payment to the specialist, and retains a commission.
- **Commission funds three things:** (1) the specialist payout, (2) payment processing fees, and (3) platform maintenance + the **free crisis path**. The paid tier cross-subsidizes the free mission — this is the platform's sustainability engine and the answer to "who pays for the free crisis help."

Specialist compensation (Year-2 decision)

Model	Description
Platform-standard rate (v1, recommended)	Every verified specialist is paid the same fixed amount per delivered session. Simple, equitable, easy to reason about.
Specialist-set rate (Prepy-style) (v2)	Each specialist sets their own per-session price; users choose on price/fit. Marketplace dynamics. Only pays off at volume.

v1 standard-rate is the ponytail choice: one number, one payout rule, no marketplace arbitrage to build. Specialist-set rates add a marketplace layer (tiers, ranking, dynamic pricing) that only pays off at volume.

Payment processing — the Venezuela complication (must flag)

Venezuela's financial system makes subscriptions and specialist payouts non-trivial:

- **Stripe does not operate in Venezuela.** Subscription billing likely requires a US/international entity (Stripe Atlas, or equivalent) or a LATAM-friendly processor (dLocal, Ebanx, MercadoPago).
- **Paying Venezuelan specialists** is the harder problem. Realistic rails in-region: Payoneer, Wise, Binance Pay (USDT), Zelle, or crypto (USDT/TRC20). Each has compliance, fee, and UX trade-offs.
- This is a **Year-2 workstream** for legal/ops, not a Year-1 task. The chosen rails affect the financial model, consumer-protection exposure, and tax reporting in both Venezuela and the processor's jurisdiction.

Trust & transparency

Introducing money changes the trust model. Rules:

- The **crisis path is and remains visibly free.** No paywall, no upsell, no "upgrade to talk now." A person in crisis never sees a price.
- Paid therapy is a clearly separate, opt-in tier — never inserted into the crisis flow.
- Pricing, session quotas, and the commission split are transparent to both users and specialists. No hidden fees.
- The landing's "personas reales" promise still holds — payment routes users to real licensed humans, not to AI.

Indicative schema (Year-2 scope, not built in Year 1)

```
subscription_plans (NEW)
  id, name, monthly_price_usd, sessions_per_month,
  status (active|retired), created_at

user_subscriptions (NEW)
  id, user_id, plan_id, status (active|cancelled|past_due),
  current_period_start, current_period_end,
  sessions_remaining, payment_processor, processor_customer_id

session_payouts (NEW)
  id, session_id, specialist_id, amount_usd,
  status (pending|paid|failed), processor_payout_id, paid_at
```

Sketched here so Year-1 schema choices (§10) don't preclude them — the cases and sessions tables are designed to attach payouts later without restructuring.

What Year 1 does to enable this

Year-1 deliverables are the *prerequisites*, not the monetization itself:

- **Verified specialist pool + scheduling + cases** (§10) → the supply side.
- **Native apps + video sessions** (§9, §5) → the delivery surface.
- **Clinical governance + legal review** (Q3) → the trust + compliance foundation that payments require.

Budget impact in Year 1

None. Subscriptions are not launched and not revenue-bearing during Year 1. The team/platform budget in §13–§14 is entirely grant/collective-funded. Payment-processing fees and specialist payouts are Year-2 lines, scoped once the model launches.

12. Team (the dominant cost)

LATAM rates (USD/month). Lean but real team. Ramps conditional on funding.

Role	Commitment	Monthly	Annual	Why
Tech lead / full-stack (current dev)	FT	$4,000–6,000	$48–72k	Owns platform + stack, mentors
Mobile engineer (React Native)	FT	$3,500–5,000	$42–60k	Owns iOS + Android (thin apps, one codebase)
Backend / platform engineer	FT (or 0.8)	$3,000–4,500	$36–54k	Owns video integration, chat (DO), reliability, data
Clinical advisor (licensed psychologist)	0.2–0.3	$800–1,500	$10–18k	Safeguarding, protocol, pro/volunteer relations
Product / UX designer	0.5–1.0	$1,500–3,000	$18–36k	UX for crisis context (low-friction is clinical)
Community / support lead	0.5	$800–1,500	$10–18k	Pro + volunteer onboarding (extends `docs/professional-communications.md`)
Legal counsel (Venezuelan lawyer, LOPDP)	retainer	$300–800	$4–10k	Volunteer charter, data protection, institutional decision
QA / test	0.5–1.0	$1,200–2,500	$14–30k	Zero tests exist today; becomes load-bearing with native + video
Grant writer / funding lead	contract (Q3–Q4)	$1,000–3,000	$6–18k	Owns funder pipeline

Team total: ~$170,000–280,000 / year.

13. Platform & services budget

Service	Cost / month	Notes
Cloudflare (Workers paid, D1 paid, R2, Analytics Engine, Email)	$5–25	Generous paid tiers; R2 grows with audio/video storage
Sentry (Team)	$26–80	Scales with native crash reporting
Apple App Store	~$8 (amortized $99/yr)	Annual fee
Google Play	~$2 (one-time $25)	Negligible
Video provider (LiveKit)	**$300–1,500**	Usage-based — the real variable cost; self-host option Q2 2027
Durable Objects (chat)	$50–300	Real-time first-aid chat
AI providers (Workers AI + LLM API)	$30–200	Pro-facing, low volume (much cheaper than user-facing triage)
Email (transactional)	$0–20	Cloudflare Email Service, currently free
Domain + DNS	~$1	psicoayudaven.com
Observability / APM (beyond Sentry)	$0–200	Optional

Platform total: ~$400–2,500 / month → ~$5,000–30,000 / year.

One-time / capex (Year 1)

Item	Cost	When
Legal review of volunteer charter + privacy (project)	$2,000–6,000	Q3
App Store / Play Store listing assets, screenshots	$500–1,500	Q4
Load testing tooling/consulting (disaster-spike readiness)	$1,000–3,000	Q4
Video integration spike (LiveKit POC)	internal	Q4

---

14. Total annual budget

Bucket	Low	High
Team (salaries)	$170,000	$280,000
Platform & services	$5,000	$30,000
One-time / capex	$4,000	$11,000
Buffer (10%)	$18,000	$32,000
TOTAL Year 1	**~$197,000**	**~$353,000**

Biggest cost lever: video adoption. If sessions take off, the $300–1,500/mo line can 3–5× by year-end. Self-hosting LiveKit (Q2 2027 if volume justifies) flattens this.

---

15. Funding strategy

Funding is a named Q3 workstream, not a side task. Target collectives (ranked by fit):

Collective	Fit	Why
Mozilla MVP (Mozillas.org)	Highest	MIT-licensed open source serving a marginalized population — textbook fit
IDB Lab / BID Lab	High	LATAM social innovation; mental health + tech eligible
Grand Challenges Canada	High	Global mental health funding; funds non-Canadian interventions
Help.NGO	High	Disaster-response tech — exact use case
Direct Relief / Americares	Medium	Health emergency NGOs with Venezuela programs
Open Society Foundations (LATAM)	Medium	Civil society + rights; longer cycle

Action: Q3 = shortlist 3–4, draft applications. Q4 = first decisions; team ramp accelerates conditional on a commitment.

Funding-contingent phasing: Q3 runs lean on existing resources (solo dev

- small legal retainer + grant writing). The Q4 team ramp accelerates only if a collective commits. If not, Q4 stays lean and we ship the volunteer flow + legal foundation on current resources, deferring native apps.

---

16. Quarterly phasing

Q3 2026 — Foundation & legal unblock (lean, pre-funding)

Lean track (always):

- Volunteer registration flow (parallel to pro registration, pointing to `/voluntariado`; cédula optional, adhesion letter captured as consent). Resolves CHANGELOG 1.23.0 explicit debt #1.
- Formal legal review kickoff (volunteer charter + privacy + AI data handling decision). Resolves debt #3.
- Institutional-registration decision (formal Venezuelan registration vs. low-profile). Resolves debt #4.
- Follow up the natural-disaster support group promised to pros on 2026-07-03 (see `docs/professional-communications.md`).
- **Funding pursuit:** shortlist collectives, draft applications.
- Design the case/session data model with the clinical advisor.
- Retention/privacy review for long-term case data.

Funded track: — (none; conditional on Q3/Q4 funding landing)

Success metric: volunteer registration ships; legal review engaged; support group has a launch date; 3+ funding applications submitted.

Q4 2026 — Volunteer model live + native scaffold + video POC

Lean track:

- Volunteer matching surface (how a help-seeker reaches a volunteer vs a verified pro).
- Landing + `/ayuda` copy rewrite (both models operational).
- Safeguarding/escalation protocol (what a volunteer does at acute risk).
- Load test for disaster spikes.

Funded track (if a collective commits):

- Hire: mobile engineer + designer + clinical advisor.
- React Native scaffold (auth + panel + directory read).

- LiveKit video integration POC.
- Durable Objects first-aid chat MVP.
- **Scheduling v1** — slot booking on D1, TZ-aware, booking confirmation (email until push exists). Migrations: `cases` + `sessions` + `slots`.

Success metric: first volunteer cohort onboarded; native app scaffold in build; LiveKit POC working end-to-end; scheduling v1 ships.

Q1 2027 — Apps in stores + video live + AI copilot v1

- iOS + Android store submissions.
- Video sessions in production (scheduled).
- Push notifications (session reminders).
- First-aid chat in production.
- **AI copilot v1:** follow-up assist + recommendations + handoff summary.
- **Case management v1** in the native app: case list, session log, treatment goals.
- Outcome metrics v1 (opt-in, aggregate-only, post-contact).
- Test infrastructure (fix the vitest/Cloudflare-plugin startup gap; tests for auth guards, `pickRandomProfessional` weighting, `isActiveNow`, per-request isolation).
- Extract `getHeaders()` to `lib/auth.ts` (the `ponytail:` ceiling, if a 4th server-fn module appears).

Success metric: apps live in both stores; first video session completed; AI copilot used by ≥1 pro; outcome metric collecting data.

Q2 2027 — Scale, sustainability, annual maturity

- Bus-factor reduction: runbooks (deploy, migration, incident), second deployer, onboarding doc.
- Governance decision (open-source contribution guide, who holds deploy keys, fiscal sponsor vs. standalone).
- AI quality / bias clinical audit (triage thresholds, false-negative rate).
- Outcome data review + first **annual impact report** (public, in Spanish; feeds back into `/acerca-de`).
- Feature-parity review vs. international peers (Crisis Text Line, 7 Cups).
- Optional Google Calendar sync if pros demand.
- Annual retrospective + replan.

Success metric: a second person can deploy; an impact report exists; the next 12 months planned from data.

17. What this plan deliberately does NOT build

- **No in-app real-time text chat as a therapy channel.** Chat is first-aid only; video is the session medium.
- **No AI chatbot as therapist / AI diagnosis / user-facing AI.** Breaks the core trust promise; AI is pro-facing only.
- **No monetization built in Year 1.** Subscriptions/specialist payouts are the post-Year-1 goal (see §11), not a Year-1 deliverable. Year 1 stays grant/collective-funded; introducing money earlier would change the trust model before the volunteer model and specialist pool are proven.
- **No multi-country expansion.** Cross-border volunteers ≠ cross-border service area. Stay Venezuela.
- **No custom video backend.** External provider (LiveKit) only.
- **No EHR.** Scheduling + cases + session logs are the minimum for continuity, not a regulated medical record system.
- **No session recording.** Live-only video; retention rule simplified.
- **No Calendly / per-seat scheduling SaaS.** Per-seat costs scale badly and data leaves the boundary.

---

18. Risks

Risk	Likelihood	Impact	Mitigation
Legal exposure of the volunteer model (unvalidated framework + safeguarding gap)	High	Existential	Q3 legal review before scaling volunteers
Funding does not land in Q3/Q4	Medium	High (delays native + video)	Funding-contingent phasing; lean Q3/Q4 keeps shipping
Video adoption spikes platform cost 3–5×	Medium	Medium	Self-host LiveKit Q2 2027 if volume justifies
D1 / Cloudflare limits under a real disaster	Medium	High	Q4 load test; verify fail-soft (the v1.21.2 infinite-skeleton pattern)

Maintainer burnout (solo → small team transition)	High	High	Q2 bus-factor work; funding enables delegation
Volunteer burnout / quality drift at scale	Medium	High	Onboarding + periodic re-engagement; clinical advisor owns
AI copilot produces a harmful recommendation that reaches a help-seeker	Low	High	Human-in-the-loop rule; pro reviews everything; Q2 bias audit

19. Open items (minimal)

Item	Owner	Resolution
AI data handling — in-region (Workers AI) vs external LLM with DPA	Legal review (Q3)	Tied to longitudinal-data privacy question
Specific funding source(s)	Funding lead	Q3 applications
Video on-demand (crisis) vs scheduled-only	Tech lead	This plan assumes scheduled-only; revisit if crisis-path video is requested

Changelog

- **2026-07-14** — v1. Initial consolidated plan. Captures decisions across team, native apps (React Native, full product), video (LiveKit), chat (Durable Objects, first-aid), AI (pro-facing copilot), scheduling + cases, retention (case-active), funding (collectives), and the backend-owns-logic architecture.